

Why Linear Types Are The Future Of Systems Programming

Aditya Siram

February 9, 2021

Introduction

Introduction

Introduction

- ATS programming language
 - ML
 - linear types
 - refinement types
 - dependant types
 - As fast as C! (“blazing fast”)
- Lots of typelevel madness
 - No optimizations
- Hongwei Xi
 - Boston University

Introduction

- Very hard!
 - Research language
 - hbox overfull with ideas
 - Tons of accidental complexity
 - Keywords everywhere . . .
 - Zero docs
- And that's OK!
 - Our job to make usable things

- Goals
 - Not evangelism!
 - Not adoption!
 - Be dissatisfied
 - Inspire your next language

Introduction

- Very difficult to present
 - Linear/dependant/refinement types, ML, C all converge
- Concrete motivating examples
 - High level handwaving
- Assuming comfort with ML like langs and basic C
- Start by taste of the ML & C side
- It'll get fairly advanced

- First from the ML side

Option datatype

- A linear Option (explanations come later ...)

```
datatype Option_vt (a:vt@type, bool) =  
  | Some_vt(a, true) of (a)  
  | None_vt(a, false)
```

Option datatype

- probably more familiar (`_vt` for viewtype)

```
datatype Option_vt =  
  | Some_vt      of (a)  
  | None_vt
```

Option datatype

- Indexed on a type-level `bool`, dependent types!

```
datatype Option_vt                                =  
  | Some_vt(a, true) of (a)  
  | None_vt(a, false)
```

Option datatype

- Sort level bool

```
datatype Option_vt                                =  
  | Some_vt(a, true) of (a)  
  | ...      ~~~~
```

Option datatype

- Parameterized on a view type, linear types!

```
datatype Option_vt (a:vt@type, bool) =  
  | ...           ~~~~~~  
  | ...
```

- All ADTs in ATS are GADTs

```
datatype Option_vt (a:vt@type, bool) =  
  | Some_vt(a, true) of (a)  
  | None_vt(a, false)
```

- A linear C array

```
absvtype arrayptr (a:vt@type, l:addr, n:int) = ptr(l)
vtypedef arrayptr (a:vt@type, n:int) =
  [l:addr] arrayptr(a, l, n)
```

Array datatype

- Just a pointer to some address, that's it

```
                                l:addr          = ptr(l)
vtypedef arrayptr              ~~~~~~
...
```

Array datatype

- Parameterized on a linear viewtype & size (should be `size_t`)

```
...  
vtypedef arrayptr (a:vt@type, n:int) =  
...           ~~~~~~
```

Array datatype

- Returns an arrayptr to an *existential* (unknown) address type

| | | |
|--------------------------------|---------------------|-----------------------|
| | <code>l:addr</code> | <code>= ptr(l)</code> |
| <code>vtypedef arrayptr</code> | <code>=</code> | |
| <code>[l:addr]</code> | | |

Array datatype

- Don't worry if this isn't clear
- Just a taste ...
- Tons type level concepts to learn!
- we'll get to some later ...

- Now from the C side!

Manual Memory Management

- What resources are leaked?

```
int main(int argc, char** argv) {  
    int* i = (int*)malloc(sizeof(int));  
    *i = 10;  
    FILE* fp = fopen("test.txt", "r");  
    return 0;  
}
```

Manual Memory Management

- Memory!

```
int main(int argc, char** argv) {  
    int* i = (int*)malloc(sizeof(int)); // <--- LEAK!!  
    *i = 10;  
    FILE* fp = fopen("test.txt", "r");  
    return 0;  
}
```

Manual Memory Management

- File descriptor

```
int main(int argc, char** argv) {  
    int* i = (int*)malloc(sizeof(int)); // <--- LEAK!!  
    *i = 10;  
    FILE* fp = fopen("test.txt", "r"); // <-- LEAK!!  
    return 0;  
}
```

Manual Memory Management

- *Equivalent* ATS program

```
implement main0 () = let
  val (pf | i) = malloc (sizeof<int>)
  val (pfset | ()) = ptr_set(pf | i, 10)
  val (pfFile | fp) = fopen("test.txt", "r")
in
  free(pfset | i);
  fclose(pfFile | fp);
end
```

Manual Memory Management

- “Client-facing” code, analogous, safe, this is why ATS is “fast”

```
implement main0 () = let
  val (      i) = malloc (sizeof<int>)
  val (      ()) = ptr_set(      i, 10)
  val (      fp) = fopen("test.txt", "r")
in
  free(      i);
  fclose(      fp);
end
```

Manual Memory Management

- `malloc` produces a linear proof `pf`, consumed by `ptr_set`

```
implement main0 () = let
  val (pf | i) = malloc (sizeof<int>)
  val (      | ()) = ptr_set(pf | i, 10)
  val (      fp) = fopen("test.txt", "r")
in
  free(      i);
  fclose(      fp);
end
```

Manual Memory Management

- `ptr_set` *produces* a proof `pfset`

```
implement main0 () = let
  val (pf | i) = malloc (sizeof<int>)
  val (pfset | ()) = ptr_set(pf | i, 10)
  val (      | fp) = fopen("test.txt", "r")
in
  free(      i);
  fclose(      fp);
end
```

Manual Memory Management

- `fopen` produces a proof of the file descriptor `pfFile`

```
implement main0 () = let
  val (pf | i) = malloc (sizeof<int>)
  val (pfset | ()) = ptr_set(pf | i, 10)
  val (pfFile | fp) = fopen("test.txt", "r")
in
  free(pfset | i);
  fclose(pfFile | fp);
end
```

Manual Memory Management

- What happens when free and fopen are deleted?

```
implement main0 () = let
  val (pf | i) = malloc (sizeof<int>)
  val (pfset | ()) = ptr_set(pf | i, 10)
  val (pfFile | fp) = fopen("test.txt", "r")
in

end
```

Manual Memory Management

- pfset is left unconsumed

```
implement main0 () = let
  val (pf | i) = malloc (sizeof<int>)
  val (pfset <---
  val (pfFile | fp) = fopen("test.txt", "r")
in

end
```

Manual Memory Management

- pfFile is left unconsumed

```
implement main0 () = let
  val (pf | i) = malloc (sizeof<int>)
  val (pfset <---
  val (pfFile <---
in

end
```

Manual Manual Management

- Consumed by free

```
implement main0 () = let
  ...
  val (pfset <----

in
  free(pfset | i); <---

end
```

Manual Memory Management

- Consumed by `fclose`, and that's it!

```
implement main0 () = let
```

```
    val (pfFile <---  
in
```

```
    fclose(pfFile | fp); <--  
end
```

Manual Memory Management

- Linear types == generalized resource tracking!
- Free to write your all your code this way!
 - safe from buffer overflows & pointer bugs
 - ... there's sugar for implicitly passing proofs around
- Reuse decades of design sensibilities (safely!)
- But you're not benefitting from Functional Programming™...

Dependant & Refinement Types

- First “big” example
 - Read a number from the user between 1 and 10
 - Allocate an array of that length
 - Fill it
 - Print it to console
 - Exit
- Doesn't seem like it but it's a LOT

Dependant & Refinement Types

- Overall structure, types simplified
- Not too far from a functional program

```
fun read_input():Option_vt(a) = ...
fun make_array (len:int n): arrayptr = ...
implement main0() = begin
  println! ("Length of array? (1-10):");
  case+ read_input<int>() of
  | ~None_vt() => println! ("Not a number!")
  | ~Some_vt(len) =>
    if (len >= 1) * (len <= 10) then
      make_array(len)
    else println! ("Bad number!")
```

- Simplified `make_array` type signature

```
fun make_array (len:int n): arrayptr = ...
```

```
...
```

```
...
```

```
...
```

- Real make_array type signature

```
fun make_array
  {n:int | n >= 1; n <= 10}
  (len:int n): [l:addr] arrayptr(int,l,n) =
  ...
```

- `len` is indexed with a refined int *sort*, `n`.

```
fun make_array
  {n:int | n >= 1; n <= 10} <-- refines it
  (len:int n): [l:addr] arrayptr(int,l,n) =
    ^^^^^
```

- Array pointer at *some* address

```
fun make_array
  {n:int | n >= 1; n <= 10}
  (len:int n): [l:addr] arrayptr(int,l,n) =
      ~~~~~~
```

- Length between 1 & 10!

```
fun make_array
  {n:int | n >= 1; n <= 10}
  (len:int n): [l:addr] arrayptr(int,l,n) =
      ~~~
```

Dependant & Refinement Types

- ... being called here

```
implement main0() =  
  ...  
  case+ ... of  
  | ...  
  | ...  
    if (len >= 1) * (len <= 10) then  
      make_array(len)  
      ~~~~~
```

Dependant & Refinement Types

- how does it know $\{n:\text{int} \mid n \geq 1; n \leq 10\}?!?$

```
implement main0() =  
  ...  
  case+ ... of  
  | ...  
  | ...  
    if (len >= 1) * (len <= 10) then  
      make_array(len)  
      ~~~~~
```

Dependant & Refinement Types

- It statically understands runtime checks!

```
implement main0() =  
  ...  
  case+ ... of  
  | ...  
  | ...  
    if (len >= 1) * (len <= 10) then  
      ~~~~~  
  ...
```

Dependant & Refinement Types

- Runtime checks discharge proofs at **compile time**.

```
implement main0() =  
  ...  
  case+ ... of  
  | ...  
  | ...  
    if (len >= 1) * (len <= 10) then  
      ~~~~~  
      ...
```

Dependant & Refinement Types

- Now anything in `make_array`'s call graph inherits the refinement

```
fun make_array
  {n:int | n >= 1; n <= 10}
  ~~~~~
  (len:int n): [l:addr] arrayptr(int,l,n) =
```

Dependant & Refinement Types

- Reading user input is actually the most interesting bit
 - It interleaves basic theorem, dependent types & runtime checks!
 - The interleaving is unique to ATS to my knowledge ...

- The old `read_input`:

```
fun read_input():Option_vt(a) = ...
```

- The actual read_input type signature:

```
fun {a:t@type} read_input():Option_vt(a) =  
  ~~~~~~                               ~~~~~~
```

Dependant & Refinement Types

- The body:

```
let
  var result: a?
  val success = fileref_load<a> (stdin_ref,result)
in
if success then
  let prval () = opt_unsome(result)
  in Some_vt(result) end
else
  let prval () = opt_unnone(result)
  in None_vt end
end
```

Dependant & Refinement Types

- Make a *stack* variable!

```
let
  var result: a? <---

in
  if success then

else

end
```

Dependant & Refinement Types

- Fill it with user input

```
let
  var result: a?
  val success = fileref_load<a> (stdin_ref,result)
in      ~~~~~~
if success then

else

end
```

Dependant & Refinement Types

- Stuff it into a Some:

```
let
  var result: a?
  val success = fileref_load<a> (stdin_ref,result)
in      ~~~~~~
if success then
  let prval () = opt_unsome(result)
  in Some_vt(result) end
else

end
```

Dependant & Refinement Types

- Hold up! result is of type $a?$, uninitialized

```
let
  var result: a? <----

in
  if success then

    in Some_vt(result) end
  else

end
```

Dependant & Refinement Types

- ...and `Option_vt(a)` needs a *not* `a`?

```
let
  var result: a? <----

in
  if success then

    in Some_vt(result) <----
  else

end
```

Dependant & Refinement Types

- The magic is happening with proof functions

```
let
  var result: a?
  val success = fileref_load <---
in
  if success then <---
    let prval () = opt_unsome(result) <---
    in Some_vt(result) end
  else

end
```

Dependant & Refinement Types

- Interleave a proof level function, erased at runtime!

```
let
  var result: a?
  val success = fileref_load
in
  if success then
    let prval () = opt_unsome(result)
    ~~~~~
  else

end
```

Dependant & Refinement Types

- Step back and look at fileref_load

```
let
  var result: a?
  val success = fileref_load <---
in
  if success then

else

end
```

- The *scary* type of `fileref_load`:

```
(FILEref, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)
```

...

- Takes a reference to stdin:

```
(FILERef, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)
~~~~~
```

- A reference (l-value) to an uninitialized stack variable:

```
(FILERef, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)
      ~~~~
```

- And returns a `bool` *indexed* with `bool`!

```
(FILERef, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)
```

~~~~~

- `success == true` indexed with a static `true`.

---

```
(FILERef, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)
```

---

~~~~~

- `failure == false` indexed with a static `false`.

```
(FILEref, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)
```

~~~~~



- The linear proof is in-place transformed ...

---

`(FILEref, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)`  
~~~~~

- ... into a tuple of an initialized a and static bool

```
(FILERef, &a? >> opt(a, b)) -<fun1> #[b:bool] bool(b)
      ~~~~~~
```

Dependant & Refinement Types

- Back to the example!

```
let
  var result: a?
  val success = fileref_load
in
  if success then

else

end
```

Dependant & Refinement Types

- success is a bool indexed with a bool

```
let
  var result: a?
  val success = fileref_load <---
in
  if success then

else

end
```

Dependant & Refinement Types

- result is a now (a,true|false)

```
let
  var result: a? <---
  val success = fileref_load
in
  if success then

else

end
```

Dependant & Refinement Types

- Now result is (a,true)!

```
let
  var result: a?
  val success = fileref_load
in
  if success then <---

else

end
```

Dependant & Refinement Types

- Now look at the *proof function* `opt_unsome`

```
let
  var result: a?
  val success = fileref_load
in
  if success then
    let prval () = opt_unsome(result) <---

  else

end
```

- The scary proof function:

```
praxi opt_unsome{a:vt@type}  
  (x: opt(a, true) >> a):<prf> void  
...
```

- It's a “proof axiom” (praxi)

```
praxi opt_unsome{a:vt@type}
```

```
~~~~~
```

```
...
```

- ... essentially a proof level assertion!

```
praxi opt_unsome{a:vt@type}
```

```
~~~~~
```

```
...
```

- In-place transforms a `opt(a, true)` into `a!`

```
praxi opt_unsome{a:vt@type}  
  (x: opt(a, true) >> a):<prf> void  
  ~~~~~
```

Dependant & Refinement Types

- So now result is a not a? !

```
let
  var result: a?
  val success = fileref_load
in
  if success then
    let prval () = opt_unsome(result)
    in Some_vt(result) end <---
  else

end
```

Dependant & Refinement Types

- `opt_unnone` does something similar!

```
let
  var result: a?
  val success = fileref_load<a> (stdin_ref,result)
in
  if success then

else
  let prval () = opt_unnone(result) <--
  in None_vt end
end
```

Dependant & Refinement Types

- Everything after `fileref_load` is purely mechanical

```
let
  var result: a?
  val success = fileref_load
in
if success then                                <--
  let prval () = opt_unsome(result) <--
  in Some_vt(result) end                      <--
else                                           <--
  let prval () = opt_unnone(result) <--
  in None_vt end                              <--
end
```

Dependant & Refinement Types

- Could all be synthesized!

```
let
  var result: a?
  val success = fileref_load
in
if success then                                <--
  let prval () = opt_unsome(result) <--
  in Some_vt(result) end                      <--
else                                           <--
  let prval () = opt_unnone(result) <--
  in None_vt end                              <--
end
```

Dependant & Refinement Types

- Taking stock ...
- Dependent types are cool
- Interleaved proof functions are a game changer
- And! ...

- Back to runtime checks!

```
fun read_input ... =  
  let  
    ...  
  in  
    if success then <---  
    else ...
```

Dependant & Refinement Types

- Back to runtime checks!

```
implement main0() =  
  ...  
  case+ ... of  
  | ... =>  
    if (len >= 1) * (len <= 10) then  
      ~~~~~
```

Proof functions

- Manipulating proof terms as 1st class citizens is a game-changer
- Can statically avoid data races!
 - Given a proof of an array of length 1 and static index i
 - Statically split it into two proofs!
 - Give each thread a sub-proof
 - Can't access other thread's array elements!
- Emulate slices!

- Proof function type signature:

```
prfun split
  {a:t@type}
  {l:addr}{n,i:nat | i <= n}
  (
    pfarr: array_v (a, l, n)
  ) : ( array_v (a, l, i),
        array_v (a, l+i*sizeof(a), n-i)
      )
...

```

Proof Functions

- `prfun ==` proof level function

`prfun split`

`...`

- Takes *proof* arguments of an array, static natural i

```
prfun split
  {a:t@type}
  {l:addr}{n,i:nat | i <= n}
  (
    pfarr: array_v (a, l, n)
  ) :
```

...

- Returns *two* proofs

```
prfun split
```

```
(
```

```
) : ( array_v (a, l, i), <--  
      array_v (a, l+i*sizeof(a), n-i) <--  
      )
```

```
...
```

- Proof of an array at l of length i

```
prfun split
```

```
(
```

```
) : ( array_v (a, l, i), <--
```

```
)
```

```
...
```

- Proof of the second section of the array!

```
prfun split
```

```
(
```

```
) : (
```

```
    array_v (a, l+i*sizeof(a), n-i) <--
```

```
)
```

```
...
```

Proof Functions

- The body

```
sif i > 0 then let
  prval (pf1, pf2arr) = array_v_uncons pfarr
  prval (pf1res1, pf1res2) =
    split{..{n-1,i-1}} (pf2arr)
in
  (array_v_cons (pf1, pf1res1), pf1res2)
end else let
  prval EQINT () = eqint_make{i,0}((*void*))
in
  (array_v_nil (), pfarr)
end
```

- There a corresponding `sif` , “static” if

```
sif i > 0 then let
```

```
in
```

```
end else let
```

```
in
```

```
end
```

Proof Functions

- Grab *proof* of the head and tail of the array

```
sif i > 0 then let
  prval (pf1, pf2arr) = array_v_uncons pfarr <--

in

end else let

in

end
```

- `array_v_uncons` is a praxi just like `opt_unsome`!

```
praxi array_v_uncons :  
{a:vt0p}{l:addr}{n:int | n > 0}  
array_v (a, l, n)  
  -<prf> (a @ l, array_v (a, l+sizeof(a), n-1))
```

Proof Functions

- Recurse with the proof of the tail and updated static counters

```
sif i > 0 then let
  prval (pf1, pf2arr) = ...
  prval (pf1res1, pf1res2) =
    split{..{n-1,i-1}} (pf2arr)
in  ~~~~~

end else let

in

end
```

Proof Functions

- Put the two sections back together!

```
sif i > 0 then let
  prval (pf1, pf2arr) = ...
  prval (pf1res1, pf1res2) =

in
  (array_v_cons (pf1, pf1res1), pf1res2) <--
end else let

in

end
```

Proof Functions

- Otherwise the first section is proof of an empty array

```
sif i > 0 then let
```

```
in
```

```
end else let
```

```
in
```

```
  (array_v_nil (), pfarr)
```

```
end
```

Proof Functions

- In a function `prval` the proofs and work in parallel!
- `thread1` and `thread2` *can not* stomp on each other!
- That's it!

```
...  
prval(pf1,pf2) = split(pfarr)  
thread1(pf1 | ...);  
thread2(pf2 | ...);  
...
```

- Tip of the iceberg!
- Proof functions means very customizable type environments
- Dependant types means much easier domain modeling
 - Skeptical “simple sum types” are sufficient
- Linear types means bullet-proof resource tracking

- All these are great ideas!
 - ATS is a great POC!
- Steadily peels back the veil
 - eg. every language designers knows proof terms
 - but keeps them internal!
 - ATS shows we're ready for them
- *The* engineering problem is UX/DX